

# Parallelization of N-body Simulation

Kai-Hsun Lu  
 National Yang Ming Chiao Tung University  
 Hsinchu, Taiwan  
 jameslu.ii14@nycu.edu.tw

Wei-Lin Wen  
 National Yang Ming Chiao Tung University  
 Hsinchu, Taiwan  
 weilin.cs14@nycu.edu.tw

Wei-Ting Yu  
 National Yang Ming Chiao Tung University  
 Hsinchu, Taiwan  
 ywt.cs14@nycu.edu.tw

## 1 Introduction / Motivation

The motivation of this project is to accelerate a physical simulation using various parallel programming techniques and to theoretically validate the expected performance improvements with practical results, so we selected the three-body simulation as a starting point and extended it to the more general N-body simulation.

The N-body problem models the evolution of a system of particles under mutual forces, however, the computational complexity of direct pairwise force evaluation is  $O(N^2)$ , making simulations with large numbers of particles computationally expensive.

To address this challenge, we aim to significantly accelerate N-body simulations and analyze the resulting performance gains with modern compute devices and various parallel programming techniques.

## 2 Statement of the Problem

**Application chosen:** N-body simulation (gravitational interactions).

**Reason for choice:**

- (1) The computation workload in N-body simulation is highly parallelizable by nature, allowing the use of different optimization approaches.
- (2) The computational complexity of  $O(N^2)$  makes it a meaningful challenge, providing a strong motivation to explore parallelism, vectorization, and efficient workload distribution.
- (3) The original implementation uses SDL for rendering, which introduces an opportunity to further investigate rendering parallelization and improve visualization performance alongside the computation.

## 3 Proposed Approaches

We will implement and compare multiple parallelization strategies for N-body simulation and introduce Tracy instrumentation methods to measure performance characteristics, identify bottlenecks, and evaluate the effectiveness of different parallel computing approaches:

- (1) Accelerations:
  - serial: original implementation.
  - Pthread pair: Directly parallelize pairwise force computations.
  - Pthread balanced: Directly parallelize but don't compute by pairwise.
  - Pthread interleaved: parallelize pairwise force computations but with interleaved index

- Mutex versions: extend all above pthread implementations to persist threads and synchronize shared data using mutexes.
  - SIMD versions: apply SIMD intrinsics to all above implementations to exploit data-level parallelism.
  - Task based parallelism
  - GPU acceleration: offload force computations to GPU using CUDA/OpenCL for massive parallelism.
- (2) step leapfrog integration:
- Identical parallelization strategies as above
- (3) rendering:
- SDL (Simple DirectMedia Layer): The original implementation uses SDL to visualize particle positions and trajectories.
  - While SDL rendering functions generally must be executed on the main thread and are not inherently parallelized, we can still parallelize certain pre-rendering tasks, such as offscreen buffer generation, where frames are computed in a memory buffer in parallel before being displayed, if time permits.

## 3.1 Block Diagram / System Architecture




Figure 1: System Architecture of the N-Body Simulation

### 3.2 Component Functions

(1) **Initialization Components:**

- **Planet Initialization:** Sets up N bodies with random positions and velocities around a center point
- **Trail Initialization:** Creates empty trail buffers for each body to track motion history
- **SDL Setup:** Initializes graphics window and surface

(2) **Physics Integration Components:**

- **step\_leapfrog():** Implements leapfrog integration algorithm with kick-drift-kick pattern: updates velocities by half-step, updates positions, recalculates accelerations, then updates velocities by another half-step
- **accelerations():** Computes gravitational forces between all bodies

(3) **Coordinate Management Components:**

- **recenter():** Calculates center of mass and shifts all bodies so the center of mass is at screen center, preventing the system from drifting off-screen

(4) **Trail Management Components:**

- **trail\_push():** Adds new position to circular trail buffer only if body moved at least MIN\_DIST from last recorded position
- **trail\_draw():** Renders motion trails with depth-based brightness

(5) **Rendering Components:**

- **fill\_circle():** Draws planets as filled circles with perspective scaling and brightness based on z-depth
- **Screen Update:** Clears screen, draws all trails and planets, updates display

(6) **Event Handling Components:**

- **SDL\_PollEvent():** Checks for quit events (window close) and handles camera control inputs for adjusting spherical coordinates  $(r, \theta, \varphi)$

(7) **Time Management Components:**

- **Frame Timer:** Uses `SDL_GetTicks()` to measure real elapsed time
- **Fixed Timestep Accumulator:** Ensures physics runs at consistent `FIXED_DT` intervals regardless of frame rate

### 3.3 Component Interactions

- (1) **Physics Update Loop:** Calculate `frame_dt` and add to accumulator and while `accumulator >= dt`, call `step_leapfrog(bodies, FIXED_DT)` which updates velocities (half-step), positions, accelerations, and velocities (half-step), then decrease accumulator
- (2) **Coordinate Management:** `recenter(bodies)` calculates center of mass and shifts all bodies to screen center
- (3) **Depth Effect:** Both `fill_circle()` and `trail_draw()` use z-coordinates to create perspective through size scaling and brightness adjustment
- (4) **Trail Recording:** For each body, `trail_push()` adds current position to trail buffer
- (5) **Rendering:** Clear screen, then for each body call `trail_draw()` and `fill_circle()`, finally `SDL_UpdateWindowSurface()`
- (6) **Tracy Integration:** `ZoneScopedN()` calls throughout enable performance profiling without affecting logic

### 3.4 Mathematical Formulations

**Leapfrog Integration (kick-drift-kick pattern):**

$$\mathbf{v}_{i+1/2} = \mathbf{v}_i + \frac{1}{2}\mathbf{a}_i\Delta t \quad (1)$$

$$\mathbf{x}_{i+1} = \mathbf{x}_i + \mathbf{v}_{i+1/2}\Delta t \quad (2)$$

$$\mathbf{a}_{i+1} = f(\mathbf{x}_{i+1}) \quad (3)$$

$$\mathbf{v}_{i+1} = \mathbf{v}_{i+1/2} + \frac{1}{2}\mathbf{a}_{i+1}\Delta t \quad (4)$$

**Center of Mass Calculation:**

$$\mathbf{r}_{COM} = \frac{\sum_{i=1}^N m_i \mathbf{r}_i}{\sum_{i=1}^N m_i} \quad (5)$$

$$\mathbf{r}'_i = \mathbf{r}_i + (\mathbf{r}_{screen} - \mathbf{r}_{COM}) \quad (6)$$

**Spherical Camera System:** The camera position is defined in spherical coordinates  $(r, \theta, \varphi)$  where  $r$  is the radial distance,  $\theta$  is the azimuthal angle, and  $\varphi$  is the polar angle.

Conversion to Cartesian coordinates:

$$x_{cam} = r \cos \theta \cos \varphi \quad (7)$$

$$y_{cam} = r \sin \theta \cos \varphi \quad (8)$$

$$z_{cam} = r \sin \varphi \quad (9)$$

**3D to 2D Projection:** Map body positions to camera view using perspective projection:

$$\mathbf{r}_{rel} = \mathbf{r}_i - \mathbf{r}_{cam} \quad (10)$$

$$x_{screen} = \frac{f \cdot x_{rel}}{z_{rel} + d} + \frac{W}{2} \quad (11)$$

$$y_{screen} = \frac{f \cdot y_{rel}}{z_{rel} + d} + \frac{H}{2} \quad (12)$$

where  $f$  is the focal length,  $d$  is the view distance, and  $(W, H)$  are screen dimensions.

### 4 Language Selection

- **C++ with pthread** for shared-memory multiprocessing, providing efficiency and fine-grained control over memory and performance with support for instrumentation tools like Tracy.
- **CUDA** for GPU acceleration, well-suited for massively parallel workloads like N-body simulations by directly mapping particles to GPU threads for concurrent force calculations.

### 5 Related Work

- Francesco, L (2025). C-projects [Source code]. GitHub. <https://github.com/mrparsing/C-Projects>
- Bartosz, Taudul (2025). tracy [Source code]. GitHub. <https://github.com/wolfpld/tracy>
- Rein van den Boomgaard (2017). Image Processing and Computer Vision, The Pinhole Camera

### 6 Expected Results

- Parallel results will accurately match the original implementation.
- Parallel implementations are expected to outperform the serial version.

- Parallel methods scale well and show speedup for different  $N$  (1K, 10K, 100K bodies).
- Performance assumptions match with Tracy measurements theoretically and mathematically.

## 7 Timetable

| Week | Task  |
|------|---|
| 3    | 1. Literature review and project planning<br>2. Implement baseline sequential N-body simulation |
| 4    | 1. Integrate Tracy profiling framework<br>2. Implement pthread parallelization                  |
| 5    | Optimize pthread implementation and profiling   |
| 6    | Implement CUDA GPU acceleration   |
| 7    | Optimize CUDA implementation and profiling  |
| 8    | Implement spherical camera system   |
| 9    | Performance comparison and bottleneck analysis  |
| 10   | Refinement and optimization of all implementations  |
| 11   | 1. Run experiments, collect performance data<br>2. Analyze results, prepare visualization       |
| 12   | Write final report and presentation   |

## 8 References

- (1) Hockney, R. W., & Eastwood, J. W. (1988). *Computer Simulation Using Particles*. CRC Press. (Leapfrog integration and N-body algorithms)
- (2) NVIDIA Corporation. (2023). *CUDA C++ Programming Guide*. <https://docs.nvidia.com/cuda/cuda-c-programming-guide/>
- (3) Tracy Profiler Documentation. (2024). <https://github.com/wolfpld/tracy/releases/latest/download/tracy.pdf>
- (4) Butcher, P., & Devlin, J. (2011). *PThreads Programming: A POSIX Standard for Better Multiprocessing*. O'Reilly Media.
- (5) SDL (Simple DirectMedia Layer) Documentation. (2024). Retrieved from <https://wiki.libsdl.org/>
- (6) Bermudez-Cameo, Jesus, Lopez-Nicolas, Gonzalo, Guerrero, Josechu. (2012). A Unified Framework for Line Extraction in Dioptric and Catadioptric Cameras.
- (7) Quinn, T., Katz, N., Stadel, J., Lake, G. (1997). Time stepping N-body simulations. <https://arxiv.org/abs/astro-ph/9710043>